


## ML Classification Project: PHDs Produced by Pakistani Universities (2010–2014)



### ✓ 1st step: Loading important .py libraries & Loading Dataset

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, ConfusionMatrixDisplay
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from xgboost import XGBClassifier

# ♦ Load Dataset
df = pd.read_csv("/content/PHDs Produced by Pakistani Universities (2010-2014).csv")
df.dropna(inplace=True)
df.head(3)
```



	S#	Institute	2010	2011	2012	2013	2014	Sector
0	1	Abdul Wali Khan University, Mardan	0	0	0	1	1	Public
1	2	Allama Iqbal Open University, Islamabad	12	13	4	4	12	Public
2	3	Air University	0	0	0	1	0	Public



Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

## ✓ 2nd Step: Preprocessing (Dataset)

```
# 📄 Step 1: Import Required Libraries
import pandas as pd
from sklearn.preprocessing import LabelEncoder, StandardScaler

# 📁 Step 2: Load Dataset
df = pd.read_csv("PHDs Produced by Pakistani Universities (2010-2014).csv")
print("✅ Loaded Dataset (First 5 Rows):\n", df.head())

# ✂️ Step 3: Drop Unnecessary Columns
df.drop(columns=["S#", "Institute"], inplace=True)
print("\n✅ After Dropping Columns ['S#", 'Institute']:\n", df.head())

# 🧑 Step 4: Handle Missing Values
missing = df.isnull().sum()
print("\n🔍 Missing Values:\n", missing)

# 🧠 Step 5: Label Encoding for Target Column
label_encoder = LabelEncoder()
df['Sector'] = label_encoder.fit_transform(df['Sector']) # Public = 1, Private = 0
print("\n✅ Encoded 'Sector' Column:\n", df['Sector'].head())

# 📊 Step 6: Feature Scaling
scaler = StandardScaler()
X = df.drop(columns=['Sector']) # Independent variables
X_scaled = scaler.fit_transform(X)
X_scaled = pd.DataFrame(X_scaled, columns=X.columns)

y = df['Sector'] # Target variable

# 📋 Output Preview
print("\n✅ Scaled Features (First 5 Rows):\n", X_scaled.head())
print("\n✅ Encoded Target (First 5 Rows):\n", y.head())
print("\n✅ Missing Values After Preprocessing:\n", df.isnull().sum())
```



```
Sector
0 Public
1 Public
2 Public
3 Public
4 Public
```

✓ After Dropping Columns ['S#', 'Institute']:

```
2010 2011 2012 2013 2014 Sector
0     0     0     0     1     1 Public
1    12    13     4     4    12 Public
2     0     0     0     1     0 Public
3    16    21    27    35    33 Public
4     0     0     1     3     0 Public
```

🔍 Missing Values:

```
2010      0
2011      0
2012      0
2013      0
2014      0
Sector     0
dtype: int64
```

✓ Encoded 'Sector' Column:

```
0     1
1     1
2     1
3     1
4     1
```

Name: Sector, dtype: int64

✓ Scaled Features (First 5 Rows):

```
2010 2011 2012 2013 2014
0 -0.419636 -0.438951 -0.467637 -0.425359 -0.456756
1  0.052454 -0.004287 -0.341859 -0.337741 -0.137622
2 -0.419636 -0.438951 -0.467637 -0.425359 -0.485768
3  0.209818  0.263199  0.381366  0.567644  0.471634
4 -0.419636 -0.438951 -0.436192 -0.366947 -0.485768
```

✓ Encoded Target (First 5 Rows):

✅ Missing Values After Preprocessing:

```
2010      0
2011      0
2012      0
2013      0
2014      0
Sector     0
dtype: int64
```

✓ 3rd Step: Apply ML Models Separately with there accuracy\_score, classification\_report, confusion\_matrix.

## (1) Logistic Regression

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
print(classification_report(y_test, y_pred_lr))
```

```
#(1) Logistic Regression
lr_model = LogisticRegression()
lr_model.fit(X_train, y_train)
y_pred_lr = lr_model.predict(X_test)
print("\n🔍 Logistic Regression Results")
print("Accuracy:", accuracy_score(y_test, y_pred_lr))
print(classification_report(y_test, y_pred_lr))
sns.heatmap(confusion_matrix(y_test, y_pred_lr), annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix - Logistic Regression")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



	precision	recall	f1-score	support
0	0.29	0.40	0.33	5
1	0.67	0.55	0.60	11
accuracy			0.50	16
macro avg	0.48	0.47	0.47	16
weighted avg	0.55	0.50	0.52	16




### Logistic Regression Results

Accuracy: 0.5

	precision	recall	f1-score	support
0	0.29	0.40	0.33	5
1	0.67	0.55	0.60	11
accuracy			0.50	16
macro avg	0.48	0.47	0.47	16
weighted avg	0.55	0.50	0.52	16

Confusion Matrix - Logistic Regression



## ✓ (2) Decision Tree


```
dt_model = DecisionTreeClassifier()
dt_model.fit(X_train, y_train)
y_pred_dt = dt_model.predict(X_test)
print("\n🔍 Decision Tree Results")
print("Accuracy:", accuracy_score(y_test, y_pred_dt))
print(classification_report(y_test, y_pred_dt))
sns.heatmap(confusion_matrix(y_test, y_pred_dt), annot=True, fmt='d', cmap='Greens')
plt.title("Confusion Matrix - Decision Tree")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



### Decision Tree Results

Accuracy: 0.8125

	precision	recall	f1-score	support
0	0.75	0.60	0.67	5
1	0.83	0.91	0.87	11
accuracy			0.81	16
macro avg	0.79	0.75	0.77	16
weighted avg	0.81	0.81	0.81	16



### ✓ (3) Random Forest

```
rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
y_pred_rf = rf_model.predict(X_test)
print("\n🔍 Random Forest Results")
print("Accuracy:", accuracy_score(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))
sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt='d', cmap='Oranges')
plt.title("Confusion Matrix - Random Forest")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```






## Random Forest Results

Accuracy: 0.5625

	precision	recall	f1-score	support
0	0.33	0.40	0.36	5
1	0.70	0.64	0.67	11
accuracy			0.56	16
macro avg	0.52	0.52	0.52	16
weighted avg	0.59	0.56	0.57	16



## ✓ (4) K-Nearest Neighbors

```
knn_model = KNeighborsClassifier()
knn_model.fit(X_train, y_train)
y_pred_knn = knn_model.predict(X_test)
print("\n🔍 KNN Results")
print("Accuracy:", accuracy_score(y_test, y_pred_knn))
print(classification_report(y_test, y_pred_knn))
sns.heatmap(confusion_matrix(y_test, y_pred_knn), annot=True, fmt='d', cmap='Purples')
plt.title("Confusion Matrix - KNN")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```




### KNN Results

Accuracy: 0.6875

	precision	recall	f1-score	support
0	0.50	0.80	0.62	5
1	0.88	0.64	0.74	11
accuracy			0.69	16
macro avg	0.69	0.72	0.68	16
weighted avg	0.76	0.69	0.70	16

Confusion Matrix - KNN



## ✓ (5) Support Vector Machine

```
svm_model = SVC(probability=True)
svm_model.fit(X_train, y_train)
y_pred_svm = svm_model.predict(X_test)
print("\n🔍 SVM Results")
print("Accuracy:", accuracy_score(y_test, y_pred_svm))
print(classification_report(y_test, y_pred_svm))
sns.heatmap(confusion_matrix(y_test, y_pred_svm), annot=True, fmt='d', cmap='Reds')
plt.title("Confusion Matrix - SVM")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```




## SVM Results

Accuracy: 0.6875

	precision	recall	f1-score	support
0	0.00	0.00	0.00	5
1	0.69	1.00	0.81	11
accuracy			0.69	16
macro avg	0.34	0.50	0.41	16
weighted avg	0.47	0.69	0.56	16

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined an
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined an
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined an
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Confusion Matrix - SVM



## ✓ (6) Naive Bayes

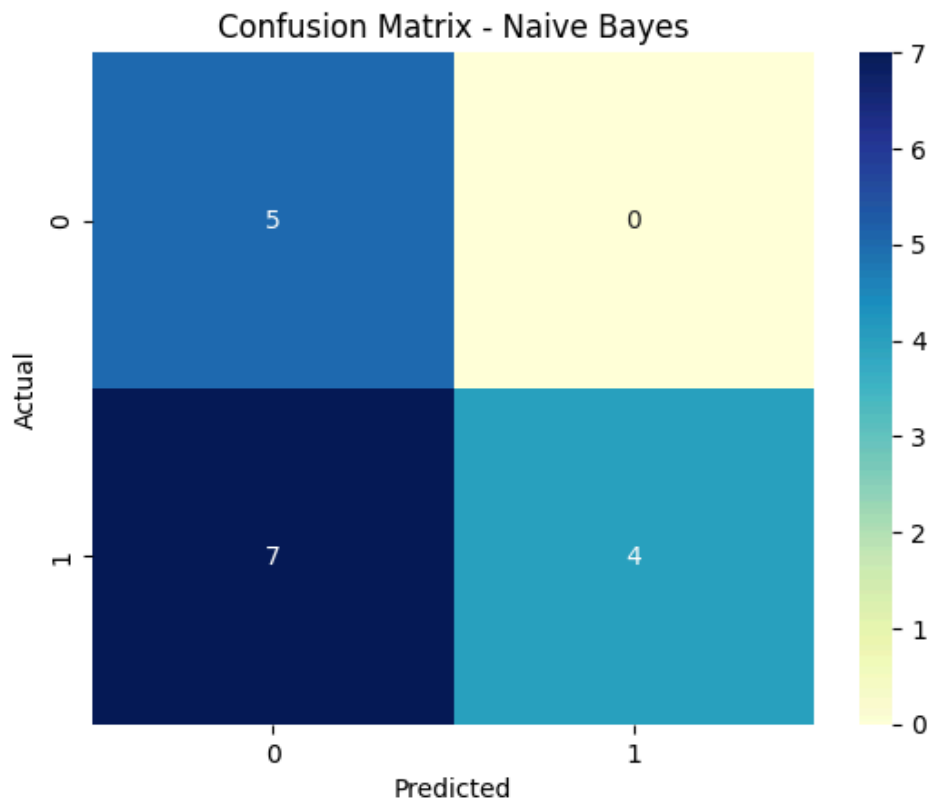
```
nb_model = GaussianNB()
nb_model.fit(X_train, y_train)
y_pred_nb = nb_model.predict(X_test)
print("\n🔍 Naive Bayes Results")
print("Accuracy:", accuracy_score(y_test, y_pred_nb))
print(classification_report(y_test, y_pred_nb))
sns.heatmap(confusion_matrix(y_test, y_pred_nb), annot=True, fmt='d', cmap='YlGnBu')
plt.title("Confusion Matrix - Naive Bayes")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



### Naive Bayes Results

Accuracy: 0.5625

	precision	recall	f1-score	support
0	0.42	1.00	0.59	5
1	1.00	0.36	0.53	11
accuracy			0.56	16
macro avg	0.71	0.68	0.56	16
weighted avg	0.82	0.56	0.55	16



- ✓ 7th Step: installing xgboost pkg and importing required libraries then performing and Evaluating Xgboost..

```
pip install xgboost
```

```
➡ Requirement already satisfied: xgboost in /usr/local/lib/python3.11/dist-packages (3.0.2)  
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.0.2)  
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.21.5)  
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from xgboost) (1.16.0)
```

```
# 📁 Required Libraries
```

```
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix  
from xgboost import XGBClassifier
```

```
# 🟡 Example Data (replace with your own dataset)
```

```
# df = pd.read_csv("your_dataset.csv")  
# X = df.drop("target", axis=1)  
# y = df["target"]
```

```
# ✅ Train-Test Split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# (7) ✅ XGBoost Classifier
```

```
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')  
xgb_model.fit(X_train, y_train)  
y_pred_xgb = xgb_model.predict(X_test)
```

```
# 📊 Evaluation
```

```
print("\n🔍 XGBoost Results")  
print("Accuracy:", accuracy_score(y_test, y_pred_xgb))  
print(classification_report(y_test, y_pred_xgb))
```

```
# 📈 Confusion Matrix
```

```
sns.heatmap(confusion_matrix(y_test, y_pred_xgb), annot=True, fmt='d', cmap='coolwarm')  
plt.title("Confusion Matrix - XGBoost")  
plt.xlabel("Predicted")  
plt.ylabel("Actual")
```



```
plt.show()
```




#### XGBoost Results

Accuracy: 0.5625

	precision	recall	f1-score	support
0	0.33	0.40	0.36	5
1	0.70	0.64	0.67	11
accuracy			0.56	16
macro avg	0.52	0.52	0.52	16
weighted avg	0.59	0.56	0.57	16

/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning: [08:14:39] WARNING: /workspace/src/learner.cc:738: Parameters: { "use\_label\_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```



TT B I <> 🔗 🖼️ “ 1/2/3 ≡ ≡ — ψ 😊 📄

#8th Step: Comparing all above applied ML-Models Performance..

## 8th Step: Comparing all above applied ML-Models Performance..

```
# 📊 Compare Models

from sklearn.metrics import precision_score, recall_score, f1_score

# 📄 Create a results list
results = []

# 🔁 Loop through all model predictions
models = {
    "Logistic Regression": y_pred_lr,
    "Decision Tree": y_pred_dt,
    "Random Forest": y_pred_rf,
    "KNN": y_pred_knn,
    "SVM": y_pred_svm,
    "Naive Bayes": y_pred_nb,
    "XGBoost": y_pred_xgb,
}

for name, preds in models.items():
    results.append({
        "Model": name,
        "Accuracy": accuracy_score(y_test, preds),
        "Precision": precision_score(y_test, preds),
        "Recall": recall_score(y_test, preds),
        "F1 Score": f1_score(y_test, preds),
    })

# 📄 Create DataFrame from results
results_df = pd.DataFrame(results)
print("✅ Model Performance Comparison:\n", results_df)

# 📊 Plot Bar Chart
plt.figure(figsize=(12, 6))
```

```
results_df.set_index("Model")[["Accuracy", "Precision", "Recall", "F1 Score"]].plot(kind="bar", figsize=(12, 6))
plt.title("📊 Model Performance Comparison")
plt.ylabel("Score")
plt.ylim(0, 1.1)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xticks(rotation=45)
plt.legend(loc="lower right")
plt.tight_layout()
plt.show()
```



## Model Performance Comparison:

	Model	Accuracy	Precision	Recall	F1 Score
0	Logistic Regression	0.5000	0.666667	0.545455	0.600000
1	Decision Tree	0.8125	0.833333	0.909091	0.869565
2	Random Forest	0.5625	0.700000	0.636364	0.666667
3	KNN	0.6875	0.875000	0.636364	0.736842
4	SVM	0.6875	0.687500	1.000000	0.814815
5	Naive Bayes	0.5625	1.000000	0.363636	0.533333
6	XGBoost	0.5625	0.700000	0.636364	0.666667

/tmp/ipython-input-1134567460.py:41: UserWarning: Glyph 128202 (\N{BAR CHART}) missing from font(s) DejaVu Sans.

plt.tight\_layout()

<Figure size 1200x600 with 0 Axes>

/usr/local/lib/python3.11/dist-packages/IPython/core/pylabtools.py:151: UserWarning: Glyph 128202 (\N{BAR CHART}) missing from font(s) DejaVu Sans.

fig.canvas.print\_figure(bytes\_io, \*\*kw)

## Model Performance Comparison

### Description of Above project

## ML Classification Project: PHDs Produced by Pakistani Universities (2010–2014)

### Project Overview

This machine learning project aims to classify whether a Pakistani university is **Public** or **Private** based on the number of PhDs it produced annually from **2010 to 2014**. The dataset was analyzed, cleaned, and used to train multiple classification models, whose performances were compared to find the best predictive approach.

### Dataset Description

- **Source:** HEC (Higher Education Commission) published records (2010–2014)

- **Columns:**

- 2010 to 2014: Number of PhDs produced in each year